

Module 13

Where Data is Stored



Exam Objective

4.3 Where Data is Stored

Objective Description

Where various types of information are stored on a Linux system.

Introduction



FSH and Processes

- A typical Linux system has thousands of files. The **Filesystem Hierarchy Standard (FHS)** provides a guideline for distributions on how to organize these files.
- The *Linux kernel* is the core of the GNU/Linux operating system. It is important to understand the role of the Linux kernel and how it both processes and provides information about the system
- Learn how to view running processes with the `ps`, `top` and other commands.
- discussion of how the system records or logs messages

Linux Kernel Processes



Kernel Processes

- A key function of the Linux kernel is to manage processes.
- The kernel accepts commands and manages processes that carry out those commands.
- The kernel gives commands access to devices like memory, disks, network interfaces, keyboards, mice, monitors and more.
- The kernel also provides access to information about active processes through a *pseudo filesystem* that is visible under the `/proc` directory. Other pseudo filesystems include `/dev` and `/sys`, which give information about hardware devices.

Pseudo filesystems are ones that appear to be real files on disk, but exist only in memory.

The /proc Directory

- The /proc directory not only contains information about processes (as name “proc” suggests), but also provides information about system hardware and current kernel configuration.
- The output shows a variety of named and numbered directories:

```
sysadmin@localhost:~$ ls /proc
```

| | | | | |
|-----|-----------|-----------|---------|---------------|
| 1 | cpuinfo | irq | modules | sys |
| 128 | crypto | kallsyms | mounts | sysrq-trigger |
| 17 | devices | kcore | mtrr | sysvipc |
| 21 | diskstats | key-users | net | thread-self |

The /proc Directory

- Some of the commands that read from `/proc` include; `top`, `free`, `mount`, `unmount`.
- There are also important regular files in the `/proc` directory such as:
 - `/proc/cmdline` - Contains information passed to kernel during boot
 - `/proc/meminfo` - Contains information about kernel memory usage
 - `/proc/modules` - Contains list of modules loaded into the kernel

Process Hierarchy

- When the kernel finishes loading during boot, it starts the *init* process and assigns it a PID of 1.
- This process then starts other system processes and assigns a PID in sequential order.
- When one process starts another process, the first process is called a *parent process*. The second process is called a *child process*.

Process Hierarchy

- Processes can be mapped into a “tree” which can be viewed with the `pstree` command.

```
sysadmin@localhost:~$ pstree
init--+-cron

      |-login---bash---pstree
      |-named---18*[{named}]
      |-rsyslogd---2*[{rsyslogd}]
      `--sshd
```

Viewing Process Snapshot

- Another way of viewing processes is with the `ps` command.
- By default, `ps` will only show running processes.
- The `ps` command can also be used with the `head` and `grep` commands to filter processes displayed:

```
sysadmin@localhost:~$ ps -e | grep firefox
6090 pts/0    00:00:07 firefox
```

Viewing Processes in Real Time

- The `top` command has a dynamic, screen-based interface that will regularly update the output of running processes.

```
sysadmin@localhost:~$ top
```

```
top - 16:58:13 up 26 days, 19:15,  1 user,  load average: 0.60, 0.74, 0.60
Tasks:   8 total,   1 running,   7 sleeping,   0 stopped,   0 zombie
Cpu(s):  6.0%us,  2.5%sy,  0.0%ni, 90.2%id,  0.0%wa,  1.1%hi,  0.2%si,  0.0%st
Mem:  32953528k total, 28126272k used,  4827256k free,    4136k buffers
Swap:          0k total,          0k used,          0k free, 22941192k cached
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-----|--------|----|----|-------|------|------|---|------|------|---------|----------|
| 1 | root | 20 | 0 | 17872 | 2892 | 2640 | S | 0 | 0.0 | 0:00.02 | init |
| 17 | syslog | 20 | 0 | 171m | 2768 | 2392 | S | 0 | 0.0 | 0:00.20 | rsyslogd |

Viewing Memory

- To view a snapshot of the memory used at that moment, use the `free` command:

```
sysadmin@localhost:~$ free
```

| | total | used | free | shared | buffers | cached |
|--------------------|----------|----------|----------|--------|---------|----------|
| Mem: | 32953528 | 26171772 | 6781756 | 0 | 4136 | 22660364 |
| -/+ buffers/cache: | | 3507272 | 29446256 | | | |
| Swap: | 0 | 0 | 0 | | | |

- The output above explained:
 - `Mem:` is the statistics for physical memory on the system
 - `-/+ buffers/cache:` is the physical memory minus memory used by the kernel
 - `Swap:` is virtual memory

Log Files

- Processes running on a system produce output that describes what the process is doing.
- Some output goes to the terminal, however other output is not seen in the terminal and gets written to files as *log messages* (or *log data*) instead.
- Some processes log data by default, while others use a daemon to log data.
 - Examples of daemons include; `syslogd`, `klogd`, `rsyslogd`, `journald`
- Log files are placed under the `/var/log` directory.

Log Files

- To view log files:
 - Use `cat` or `less` command
 - Use `journalctl` command
- Log files are *rotated*, meaning older log files are renamed and replaced with newer log files.
- Most log files contain text, which can be viewed safely with many tools. Other files such as the `/var/log/btmp` and `/var/log/wtmp` files contain binary. Use the `file` command to view binary log files.

Kernel Messages

- Kernel messages can be found in the following files:
 - `/var/log/dmesg` - contains the kernel messages that were produced during system startup.
 - `/var/log/messages` - will contain kernel messages that are produced as the system is running.
- To view messages generated by the kernel, use the `dmesg` command. To filter the output, use a pipe with the `less` or `grep` command:

```
sysadmin@localhost:~$ dmesg | grep -i usb
usbcore: registered new interface driver usbfs
usbcore: registered new interface driver hub
usbcore: registered new device driver usb
```


Filesystem Hierarchy Standard



Filesystem Hierarchy Standard

- Filesystem Hierarchy Standard (FHS) is a set of standards supported by the Linux Foundation.
- FHS categorizes system directories as:
 - Shareable / Not shareable
 - Static / Variable
- The FHS standard defines four hierarchies of directories used in organizing the files of the filesystem:
 - Top-level hierarchy: /
 - Second-level hierarchy: `/usr`
 - Third-level hierarchy: `/usr/local`
 - Fourth-level hierarchy: `/var`

Organization Within the Filesystem Hierarchy

- **User and Home Directories:** The `/home` directory will typically have a directory underneath it for each user account.
- **Binary Directories:** Contains the programs that users and administrators execute to start processes or applications running on the system.
 - Includes `/bin`, `/usr/bin`, `/usr/local/bin` and other non-user specific directories.
- **Root Restricted Binaries:** the `sbin` directories are primarily intended to be used by the system administrator (the root user) and include:
 - `/sbin`, `/usr/sbin`, and `/usr/local/sbin`

Organization Within the Filesystem Hierarchy

- **Software Application Directories:**

- Microsoft Windows - Applications files are installed in a single subdirectory under the `C:\Program Files` directory.
- Linux - Applications may have files in multiple directories spread out throughout the Linux filesystem.
- To view list of application files, use `dpkg -L packagename` (Debian) and `rpm -ql packagename` (Red Hat).

- **Library Directories:** Files which contain code that is shared between multiple programs.

- Commonly use file extension of `.so`
- Examples include: `/lib`, `/lib64`, `/usr/lib`, `/usr/lib64`, `/usr/local/lib`

Organization Within the Filesystem Hierarchy

- **Variable Data Directories:** The `/var` directory and many of its subdirectories can contain data that will change frequently.
 - Examples include: `/var/mail`, `/var/spool/mail`, `/var/spool/cups`